

Programming Tools For Image Processing

Traditional Image Processing
Neural Networks

Today's Lecture


- Oriented towards the programming assignments
 - Image processing with numpy
 - Convolutional operations
 - Backpropagation
 - Neural Networks in Pytorch

Representation of Images

```
1  im = np.array([
2      [255, 255, 255],
3      [255, 0, 255],
4      [255, 255, 255]
5  ])
6  plt.imshow(im, cmap="gray")
```

Can be represented in the range 0-255 (8bit unsigned int)
Shape: (Image Height, Image Width, #color channels)


Representation of Images



```
1  im = np.array([  
2      [255, 255, 255],  
3      [255, 0, 255],  
4      [255, 255, 255]  
5  ])  
6  plt.imshow(im, cmap="gray")
```

Can be represented in the range 0-255 (8bit unsigned int)
Shape: (Image Height, Image Width, #color channels)


Representation of Images



```
1  im = np.array([
2      [1., 1., 1.],
3      [1., 0., 1.],
4      [1., 1., 1.]
5  ])
6  plt.imshow(im, cmap="gray")
```


Can be represented in the range 0-1 (float)
Shape: (Image Height, Image Width, #color channels)

Working with Numpy




```
1  # We can make this image  
2  # into RGB (3 dimensions)  
3  im = im.reshape(3, 3, 1)  
4  im = np.tile(im, [1, 1, 3])  
5  print(im.shape)
```

Working with Numpy



```
1  # We can set the color
2  # "blue" to 0
3  im[:, :, 2] = 0
4  plt.imshow(im)
```

Working with Numpy



```
1 # We can set the color
2 # "blue" to 0
3 im[:, :, 2] = 0
4 plt.imshow(im)
```


Working with Numpy

```
1 import skimage
2 chelsea = skimage.data.chelsea()
3 plt.imshow(chelsea)
4 print(f"Image has shape: {chelsea.shape}, ",
5       f"with dtype={chelsea.dtype}, ",
6       f"min value={chelsea.min()}, ",
7       f"max value={chelsea.max()}")
```

Image has shape: (300, 451, 3),
with dtype=uint8,
min value=0, max value=231

Working with Numpy



```
1 import skimage
2 chelsea = skimage.data.chelsea()
3 plt.imshow(chelsea)
4 print(f"Image has shape: {chelsea.shape}, ",
5       f"with dtype={chelsea.dtype}, ",
6       f"min value={chelsea.min()}, ",
7       f"max value={chelsea.max()}")
```

Image has shape: (300, 451, 3),
with dtype=uint8,
min value=0, max value=231

Working with Numpy



```
1  # Lets put a green box in the middle
2  chelsea_1 = chelsea.copy()
3  chelsea_1[100:150, 200:300, :] = 0
4  chelsea_1[100:150, 200:300, 1] = 255
5  plt.imshow(chelsea_1)
```

Row, Column, Color


Working with Numpy




```
1  # Lets put a green box in the middle
2  chelsea_1 = chelsea.copy()
3  chelsea_1[100:150, 200:300, :] = 0
4  chelsea_1[100:150, 200:300, 1] = 255
5  plt.imshow(chelsea_1)
```

Row, Column, Color

Convolutional operation



Convolutional operation




Convolutional operation

1. Place the kernel in the top left corner
2. Apply the kernel and compute the result (single number)
3. Slide the kernel to the right by 1 pixel

f = image, h = convolutional kernel

$$(f * h)(x, y) = \sum_{i=-\infty}^{\infty} \sum_{j=-\infty}^{\infty} f(i, j)(h(x - i, y - j))$$




Convolutional operation

1. Place the kernel in the top left corner
2. Apply the kernel and compute the result (single number)
3. Slide the kernel to the right by 1 pixel

f = image, h = convolutional kernel

$$(f * h)(x, y) = \sum_{i=-\infty}^{\infty} \sum_{j=-\infty}^{\infty} f(i, j)(h(x - i, y - j))$$



Convolution Operation Usage

We can find edges by using Sobel:



Convolution Operation Usage


We can find edges by using Sobel (Vertical edges):



1	0	-1
2	0	-2
1	0	-1

Convolution Operation Usage

We can find edges by using Sobel (Vertical edges):



1	0	-1
2	0	-2
1	0	-1

Convolution Operation Usage


We can find edges by using Sobel (horizontal):



1	2	1
0	0	0
-1	-2	-1

Convolution Operation Usage



We can find edges by using Sobel (horizontal):



1	2	1
0	0	0
-1	-2	-1

Convolution Operation Usage

Can combine the horizontal and vertical edges



$$G = \sqrt{G_x^2 + G_y^2}$$

Convolution Operation Usage

Can combine the vertical and horizontal edges



$$G = \sqrt{G_x^2 + G_y^2}$$

Convolution Example

Convolve the 3x5 image with the 3x3 Kernel

1	2	3	4	5
7	8	9	10	11
12	13	14	15	16

0	1	0
1	0	1
0	2	0

$$(f * h)(x, y) = \sum_{i=-\infty}^{\infty} \sum_{j=-\infty}^{\infty} f(i, j)(h(x - i, y - j))$$

Convolution Example

Convolve the 3x5 image with the 3x3 Kernel

Instead, we can rotate the kernel and perform correlation.

1	2	3	4	5
7	8	9	10	11
12	13	14	15	16

0	2	0
1	0	1
0	1	0

(Kernel Rotated 180°)

Convolution Example

Convolve the 3x5 image with the 3x3 Kernel

Instead, we can rotate the kernel and perform correlation.

1	2	3	4	5
7	8	9	10	11
12	13	14	15	16

0	2	0
1	0	1
0	1	0

(Kernel Rotated 180°)



Start in the top left (legal) corner

Convolution Example

Convolve the 3x5 image with the 3x3 Kernel

Instead, we can rotate the kernel and perform correlation.

1	2	3	4	5
7	8	9	10	11
12	13	14	15	16

0	2	0
1	0	1
0	1	0

(Kernel Rotated 180°)

33

$$2*2 + 7*1 + 1*9 + 1*13 = 33$$

Start in the top left (legal) corner

Convolution Example

Convolve the 3x5 image with the 3x3 Kernel

Instead, we can rotate the kernel and perform correlation.

1	2	3	4	5
7	8	9	10	11
12	13	14	15	16

0	2	0
1	0	1
0	1	0

(Kernel Rotated 180°)

33	38	
----	----	--

Stride along the image (horizontally and vertically)

Convolution Example

Convolve the 3x5 image with the 3x3 Kernel

Instead, we can rotate the kernel and perform correlation.

1	2	3	4	5
7	8	9	10	11
12	13	14	15	16

0	2	0
1	0	1
0	1	0

(Kernel Rotated 180°)

33	38	43
----	----	----

Stride along the image (horizontally and vertically)

Convolution Example - With Padding

Convolve the 3x5 image with the 3x3 Kernel

0	0	0	0	0	0	0
0	1	2	3	4	5	0
0	7	8	9	10	11	0
0	12	13	14	15	16	0
0	0	0	0	0	0	0

0	2	0
1	0	1
0	1	0

(Kernel Rotated 180°)

Convolution Example - With Padding

Convolve the 3x5 image with the 3x3 Kernel

0	0	0	0	0	0	0
0	1	2	3	4	5	0
0	7	8	9	10	11	0
0	12	13	14	15	16	0
0	0	0	0	0	0	0

0	2	0
1	0	1
0	1	0

(Kernel Rotated 180°)

To keep the original shape, we pad with some value
Most common:

- Reflection Padding
- Zero Padding

Convolution Example - With Padding

Convolve the 3x5 image with the 3x3 Kernel

0	0	0	0	0	0	0
0	1	2	3	4	5	0
0	7	8	9	10	11	0
0	12	13	14	15	16	0
0	0	0	0	0	0	0

0	2	0
1	0	1
0	1	0

(Kernel Rotated 180°)

9	12	15	18	15
22	33	38	43	36
27	42	46	50	37

To keep the original shape, we pad with some value

Most common:

- Reflection Padding
- Zero Padding

Classification with Neural Networks

We are going to build a NN to classify digits (0-10)

We will use the MNIST database:

- 60,000 images in training set
- 10,000 images in test set



MNIST Neural Network


Each image has shape 28x28x1



Want to predict the confidence
(0-1) for each digit

MNIST Neural Network

Each image has shape 28x28x1




Want to predict the confidence
(0-1) for each digit

MNIST Neural Network

Each image has shape 28x28x1



Want to predict the confidence (0-1) for each digit




MNIST Neural Network

Each image has shape 28x28x1



Want to predict the confidence (0-1) for each digit




MNIST Neural Network

Each image has shape 28x28x1




Want to predict the confidence (0-1) for each digit



MNIST Neural Network

We will have:


- y : 10 outputs
- x : 784 inputs



MNIST Neural Network

We will have:

- y : 10 outputs
- x : 784 inputs
- w : [784, 10]
- bias: 10 (per output)



MNIST Neural Network - In Numpy

We will have:

- y: 10 outputs
- x: 784 inputs
- w: [784, 10]
- bias: 10 (per output)

```
1  # Simple neural network
2  def forward(x, w):
3      z = x.dot(w)
4      a = softmax(z)
5      return a
```

MNIST Neural Network - In Numpy

Our neural networks outputs confidence scores



→ [0.9, 0.01, 0.01, 0.01, 0.01, 0.01, 0.01, 0.01, 0.02, 0.01]

Digit "0"

Digit "7"

Training a Neural Network

We have to define a cost function:

$$C = - \sum_n^N \sum_{k=1}^C \hat{y}_k^n \ln(y_k^n)$$

Cross entropy loss

Measures “how good” our classification is over N training examples


Training a Neural Network

We have to define a cost function:

$$E = - \sum_n \sum_{k=1}^C \hat{y}_k^n \ln(y_k^n)$$


Cross entropy loss

Measures “how good” our classification is over N training examples



Gradient descent

- The building block of all neural networks

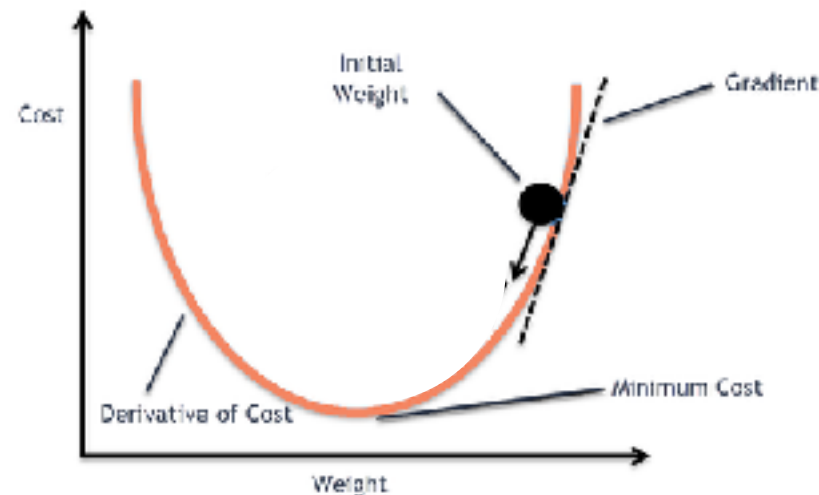


Gradient descent

- The building block of all neural networks
- Minimize the objective function

$$w_{t+1} = w_t - \alpha \frac{\partial C^n(w)}{\partial \theta}$$

- alpha: learning rate




Gradient descent

- The building block of all neural networks
- Minimize the objective function

$$w_{t+1} = w_t - \alpha \frac{\partial C^n(w)}{\partial \theta}$$

- alpha: learning rate




Backpropagation example

Let's say:

$$C = (y - \hat{y})^3$$

And:

$$\hat{y} = 5$$



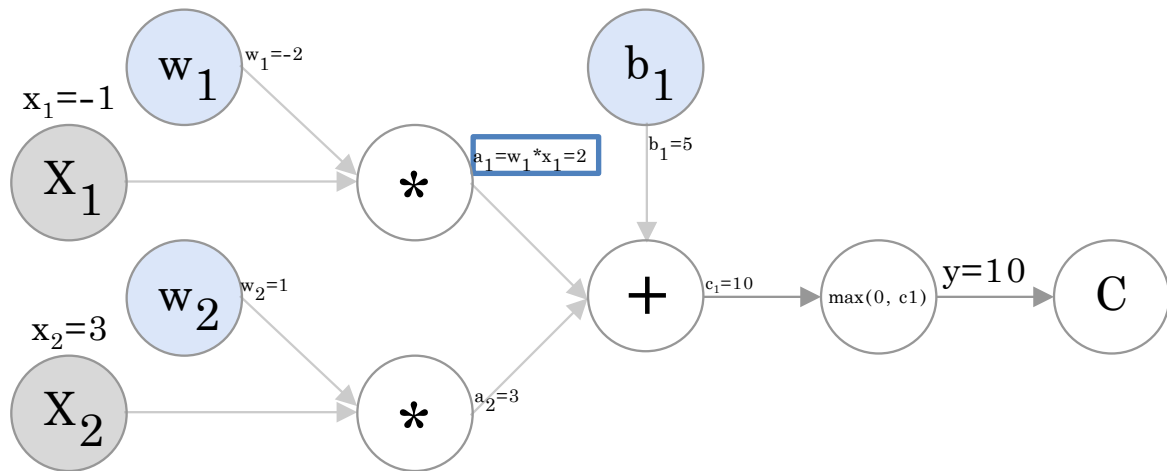
Backpropagation example

Let's say:

$$C = (y - \hat{y})^3$$

And:

$$\hat{y} = 5$$



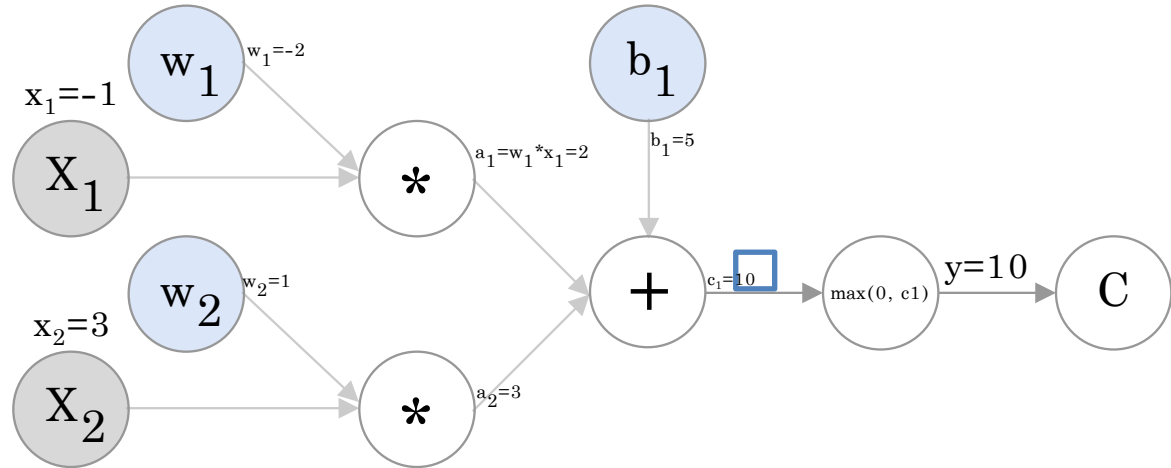
Backpropagation example

Let's say:

$$C = (y - \hat{y})^3$$

And:

$$\hat{y} = 5$$



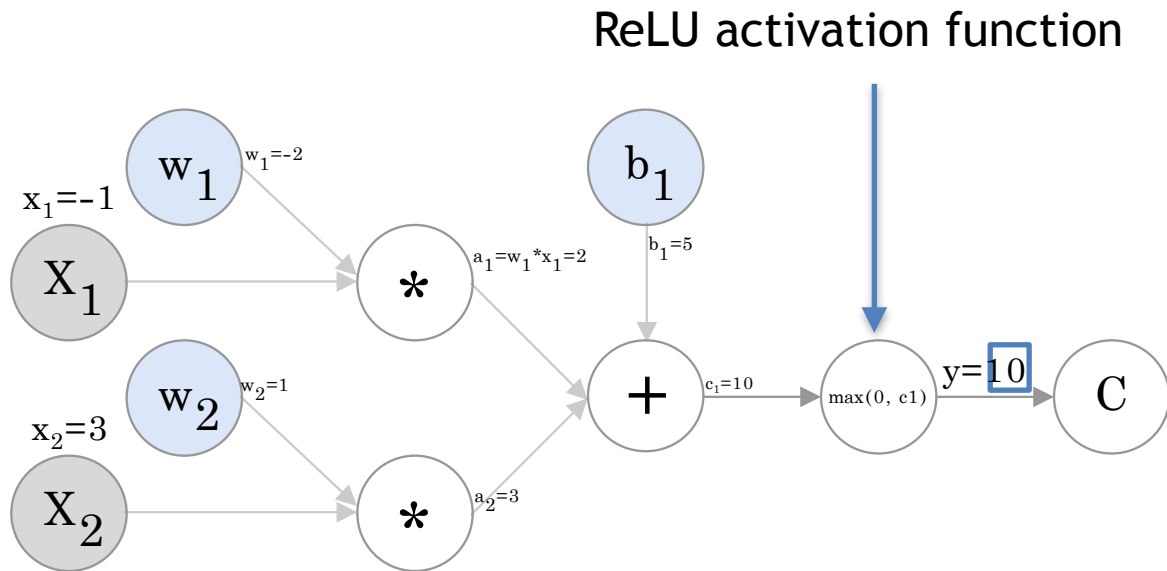
Backpropagation example

Let's say:

$$C = (y - \hat{y})^3$$

And:

$$\hat{y} = 5$$



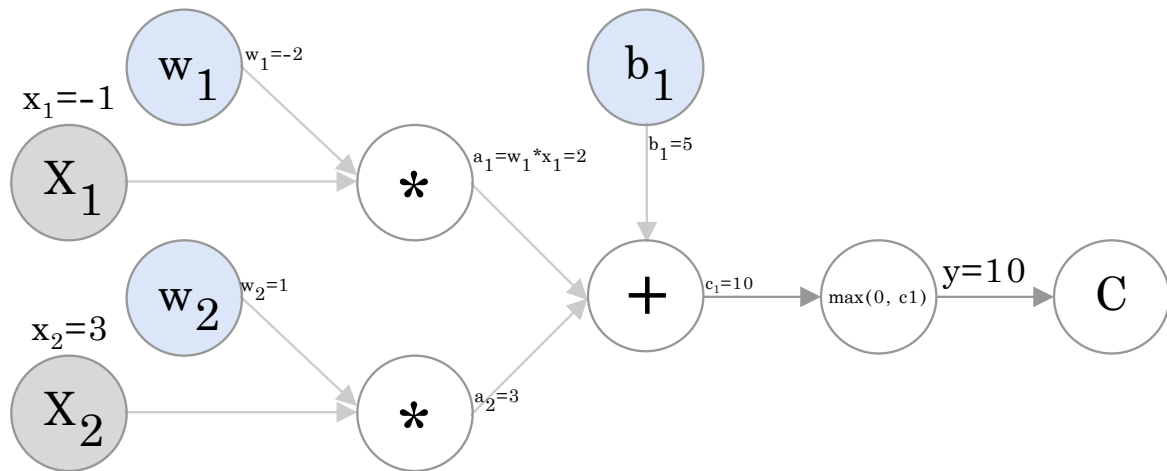
Backpropagation example

Let's say:

$$C = (y - \hat{y})^3$$

And:

$$\hat{y} = 5$$

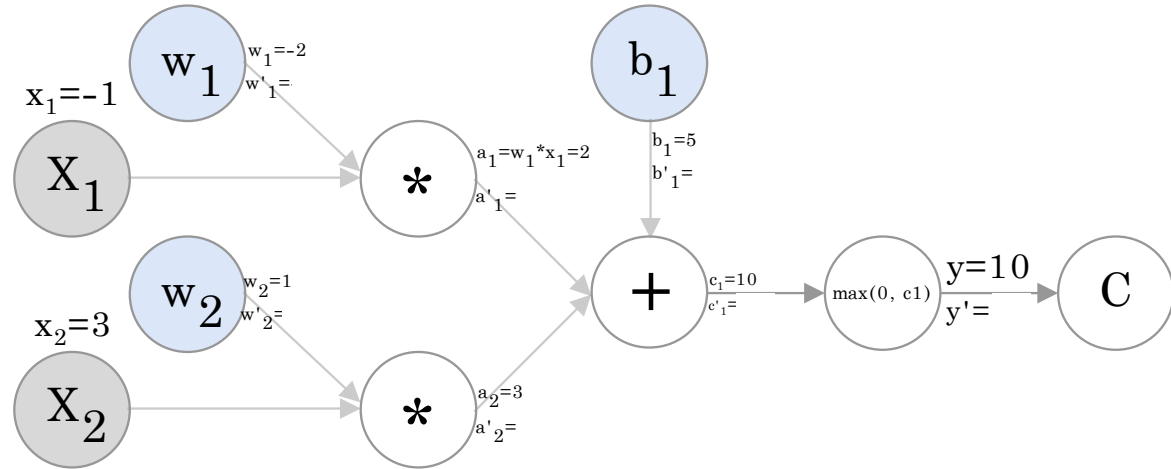


$$C = 125$$

Backpropagation example

We know that

$$1. \frac{\partial C}{\partial y} = 3 * (y - \hat{y})^2$$

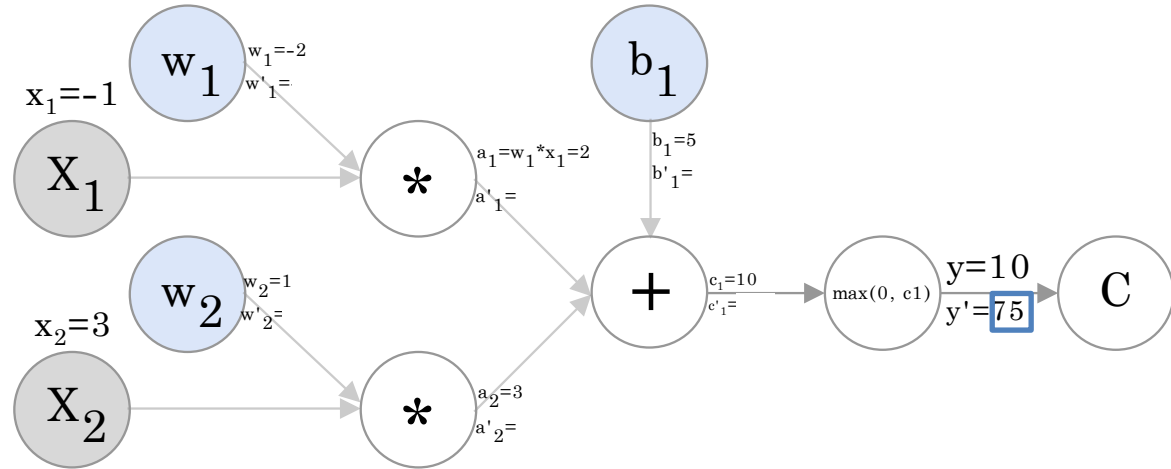


$$C = (y - \hat{y})^3, \hat{y} = 5$$

Backpropagation example

We know that

$$1. \frac{\partial C}{\partial y} = 3 * (y - \hat{y})^2 = 75$$



$$C = (y - \hat{y})^3, \hat{y} = 5$$

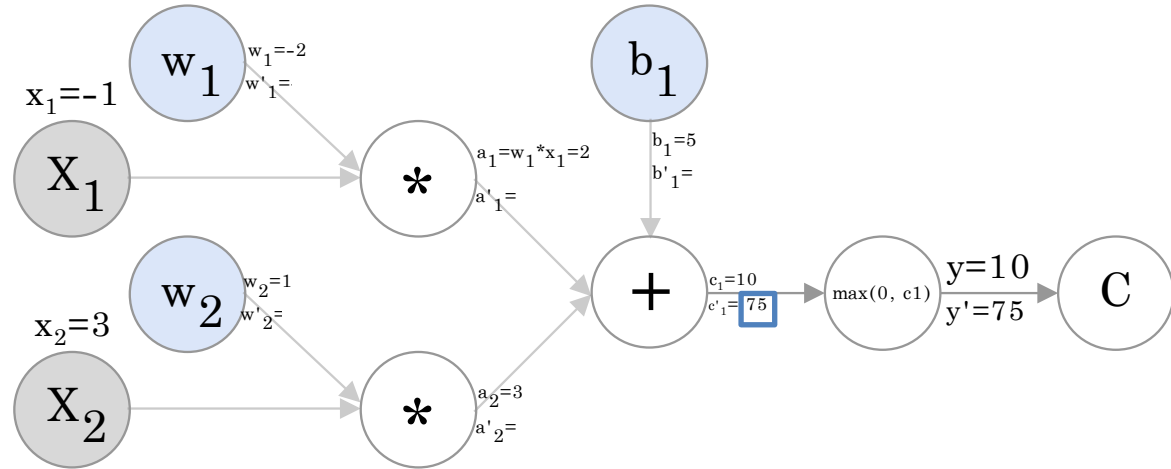
Backpropagation example

We know that

$$1. \frac{\partial C}{\partial y} = 3 * (y - \hat{y})^2 = 75$$

Then by using chain rule:

$$2. \frac{\partial C}{\partial c_1} = \frac{\partial C}{\partial y} \frac{\partial y}{\partial c_1} = 75 * 1 = 75$$



$$C = (y - \hat{y})^3, \hat{y} = 5$$

Backpropagation example

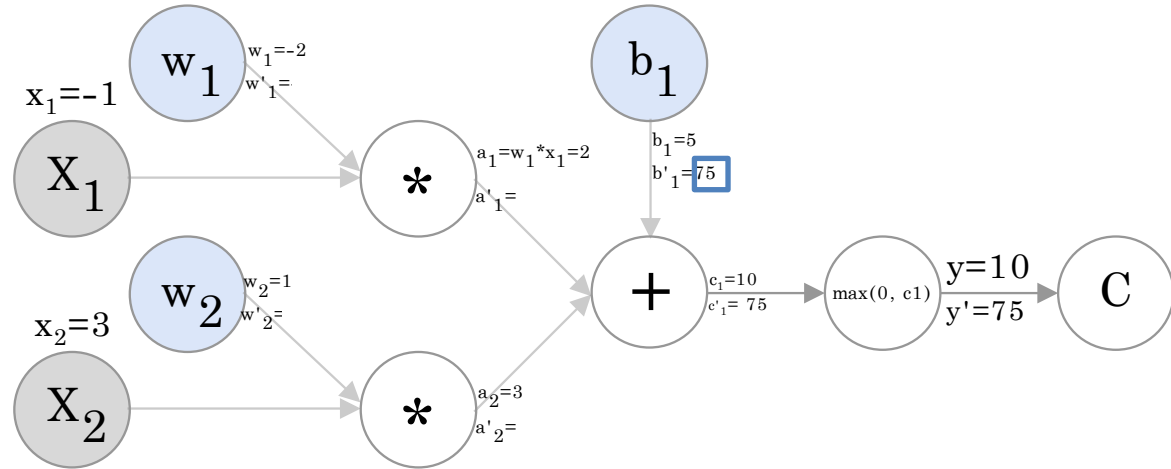
We know that

$$1. \frac{\partial C}{\partial y} = 3 * (y - \hat{y})^2 = 75$$

Then by using chain rule:

$$2. \frac{\partial C}{\partial c_1} = \frac{\partial C}{\partial y} \frac{\partial y}{\partial c_1} = 75 * 1 = 75$$

$$3. \frac{\partial C}{\partial b_1} = \frac{\partial C}{\partial c_1} \frac{\partial c_1}{\partial b_1} = 75 * 1 = 75$$



$$C = (y - \hat{y})^3, \hat{y} = 5$$

Backpropagation example

We know that

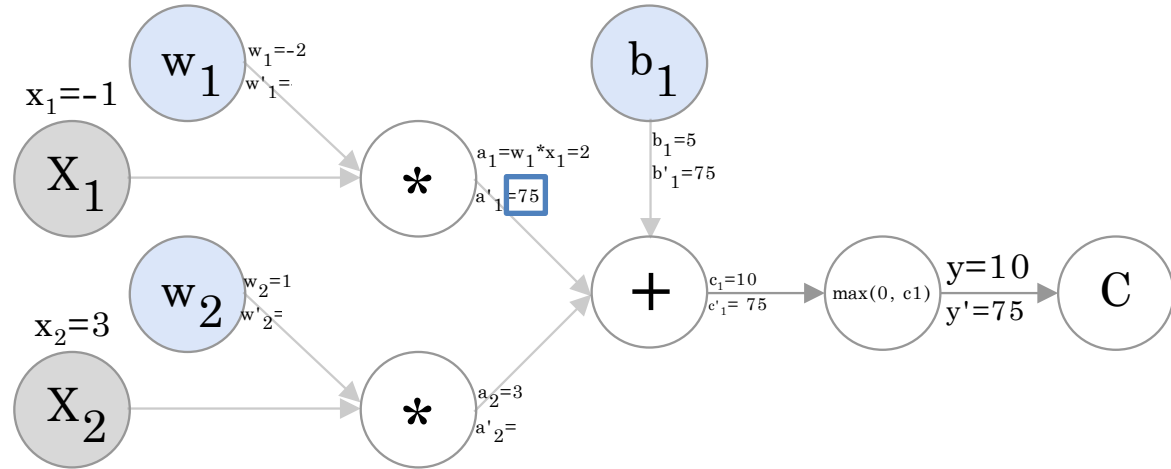
$$1. \frac{\partial C}{\partial y} = 3 * (y - \hat{y})^2 = 75$$

Then by using chain rule:

$$2. \frac{\partial C}{\partial c_1} = \frac{\partial C}{\partial y} \frac{\partial y}{\partial c_1} = 75 * 1 = 75$$

$$3. \frac{\partial C}{\partial b_1} = \frac{\partial C}{\partial c_1} \frac{\partial c_1}{\partial b_1} = 75 * 1 = 75$$

$$4. \frac{\partial C}{\partial a_1} = \frac{\partial C}{\partial c_1} \frac{\partial c_1}{\partial a_1} = 75 * 1 = 75$$



$$C = (y - \hat{y})^3, \hat{y} = 5$$

Backpropagation example

We know that

$$1. \frac{\partial C}{\partial y} = 3 * (y - \hat{y})^2 = 75$$

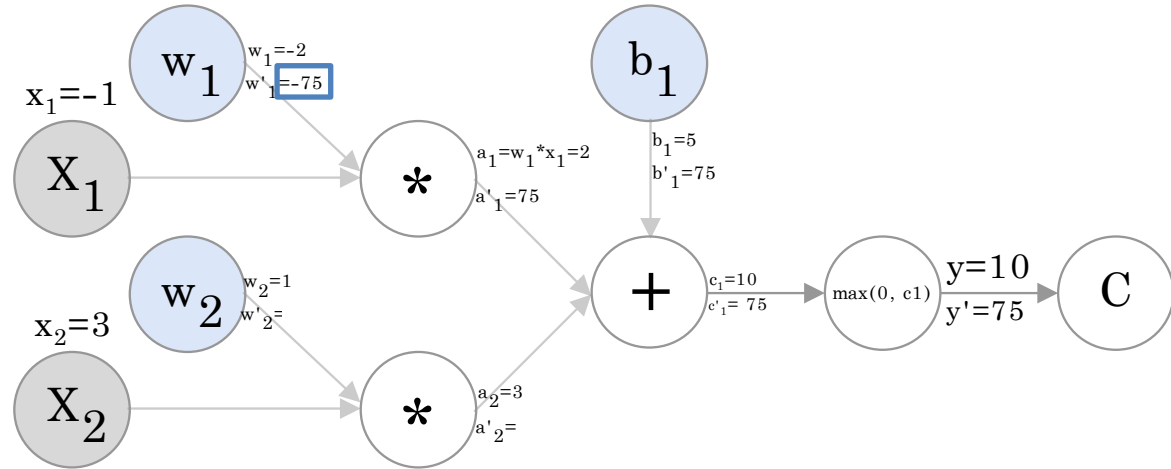
Then by using chain rule:

$$2. \frac{\partial C}{\partial c_1} = \frac{\partial C}{\partial y} \frac{\partial y}{\partial c_1} = 75 * 1 = 75$$

$$3. \frac{\partial C}{\partial b_1} = \frac{\partial C}{\partial c_1} \frac{\partial c_1}{\partial b_1} = 75 * 1 = 75$$

$$4. \frac{\partial C}{\partial a_1} = \frac{\partial C}{\partial c_1} \frac{\partial c_1}{\partial a_1} = 75 * 1 = 75$$

$$5. \frac{\partial C}{\partial w_1} = \frac{\partial C}{\partial a_1} \frac{\partial a_1}{\partial w_1} = 75 x_1 = -75$$



$$C = (y - \hat{y})^3, \hat{y} = 5$$

Backpropagation example

We know that

$$1. \frac{\partial C}{\partial y} = 3 * (y - \hat{y})^2 = 75$$

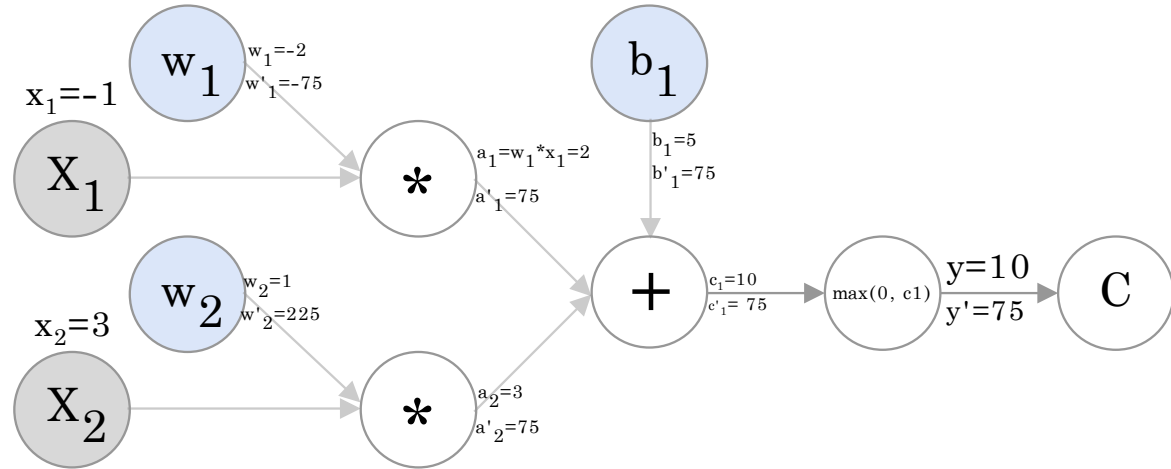
Then by using chain rule:

$$2. \frac{\partial C}{\partial c_1} = \frac{\partial C}{\partial y} \frac{\partial y}{\partial c_1} = 75 * 1 = 75$$

$$3. \frac{\partial C}{\partial b_1} = \frac{\partial C}{\partial c_1} \frac{\partial c_1}{\partial b_1} = 75 * 1 = 75$$

$$4. \frac{\partial C}{\partial a_1} = \frac{\partial C}{\partial c_1} \frac{\partial c_1}{\partial a_1} = 75 * 1 = 75$$

$$5. \frac{\partial C}{\partial w_1} = \frac{\partial C}{\partial a_1} \frac{\partial a_1}{\partial w_1} = 75 x_1 = -75$$




$$C = (y - \hat{y})^3, \hat{y} = 5$$

Gradient descent

Our update rule ($\alpha = .01$):

$$w_1 = w_1 - \alpha \frac{\partial C}{\partial w_1}$$



Gradient descent

Our update rule ($\alpha = .01$):


$$w_1 = w_1 - \alpha \frac{\partial C}{\partial w_1}$$

We know:

$$\frac{\partial C}{\partial w_1} = -75, \text{ and } w_1 = -2$$

Then,

$$w_1 = -2 - 0.01 \cdot (-75) = -1.25$$

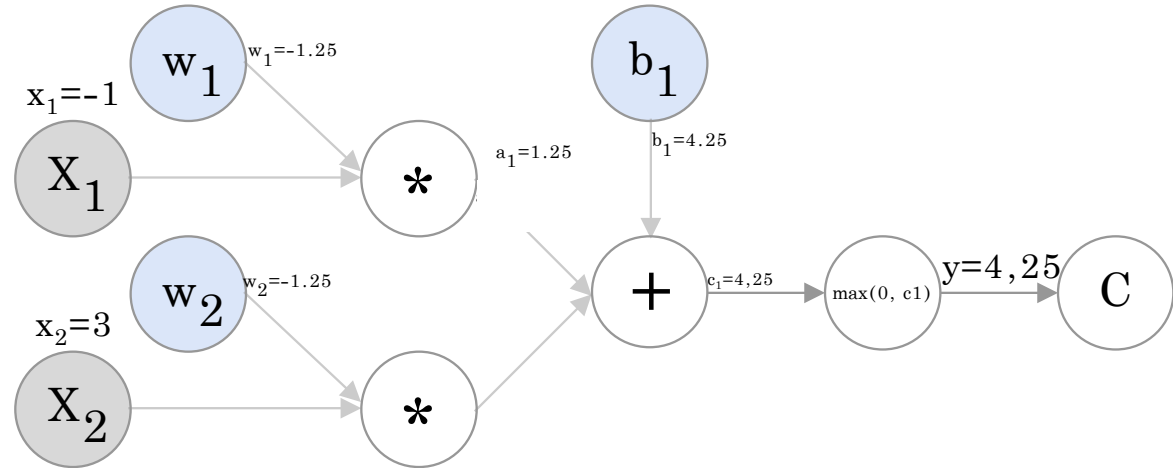


Gradient descent - Did we get closer?

Yes!

New prediction:

$$y = 4.25$$



Neural Network/Machine Learning Concepts

Hyperparameters:

- Parameters we set before training.
 - Learning rate
 - Batch size

Neural Network/Machine Learning Concepts

Hyperparameters:

- Parameters we set before training.
 - Learning rate
 - Batch size

Minibatch:

- Instead of updating weights on a single training example, we take the average over a minibatch

Neural Network/Machine Learning Concepts

Hyperparameters:

- Parameters we set before training.
 - Learning rate
 - Batch size

Minibatch:


- Instead of updating weights on a single training example, we take the average over a minibatch

Datasets:

- Train: Used only for training
- Validation: Used to validate model / tune hyperparameters
- Test: Final evaluation of model (should not be used frequently)

Overfitting

- Model memorise training points
- Does not generalise to the underlying function



Neural Network/Machine Learning Concepts

Hyperparameters:

- Parameters we set before training.
 - Learning rate
 - Batch size

Minibatch:

- Instead of updating weights on a single training example, we take the average over a minibatch

Datasets:

- Train: Used only for training
- Validation: Used to validate model / tune hyperparameters
- Test: Final evaluation of model (should not be used frequently)

Gradient Descent - In Numpy

```
1 def gradient_decent(X, outputs, targets, weights):
2     N = X.shape[0]
3     for i in range(weights.shape[0]):
4         dw_i = - 2 * (targets - outputs) * X[:, i:i+1]
5
6         dw_i = dw_i.mean(axis=0)
7         weights[i] = weights[i] - learning_rate * dw_i
8     return weights
9
```

Gradient Descent - In Numpy

For every weight

```
1 def gradient_decent(X, outputs, targets, weights):
2     N = X.shape[0]
3     for i in range(weights.shape[0]):
4         dw_i = - 2 * (targets - outputs) * X[:, i:i+1]
5
6         dw_i = dw_i.mean(axis=0)
7         weights[i] = weights[i] - learning_rate * dw_i
8     return weights
9
```

Gradient Descent - In Numpy

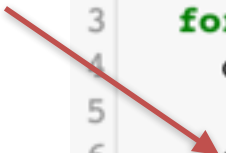
Compute gradient

```
1 def gradient_decent(X, outputs, targets, weights):
2     N = X.shape[0]
3     for i in range(weights.shape[0]):
4         dw_i = - 2 * (targets - outputs) * X[:, i:i+1]
5
6         dw_i = dw_i.mean(axis=0)
7         weights[i] = weights[i] - learning_rate * dw_i
8     return weights
9
```

Gradient Descent - In Numpy

Find mean over all training examples

```
1 def gradient_decent(X, outputs, targets, weights):
2     N = X.shape[0]
3     for i in range(weights.shape[0]):
4         dw_i = - 2 * (targets - outputs) * X[:, i:i+1]
5
6         dw_i = dw_i.mean(axis=0)
7         weights[i] = weights[i] - learning_rate * dw_i
8     return weights
9
```



Gradient Descent - In Numpy

Perform update



```
1 def gradient_decent(X, outputs, targets, weights):
2     N = X.shape[0]
3     for i in range(weights.shape[0]):
4         dw_i = - 2 * (targets - outputs) * X[:, i:i+1]
5
6         dw_i = dw_i.mean(axis=0)
7         weights[i] = weights[i] - learning_rate * dw_i
8     return weights
9
```


Gradient Descent - In Numpy

This is a lot of code for a single layer!

And it's extremely slow

```
1 def gradient_decent(X, outputs, targets, weights):
2     N = X.shape[0]
3     for i in range(weights.shape[0]):
4         dw_i = - 2 * (targets - outputs) * X[:, i:i+1]
5
6         dw_i = dw_i.mean(axis=0)
7         weights[i] = weights[i] - learning_rate * dw_i
8     return weights
9
```

Instead, use a framework

Why?

- Quickly implement and test ideas
- Automatically compute gradients
- Run it efficient

Frameworks

Caffe2

(Facebook/
UC Berkeley)

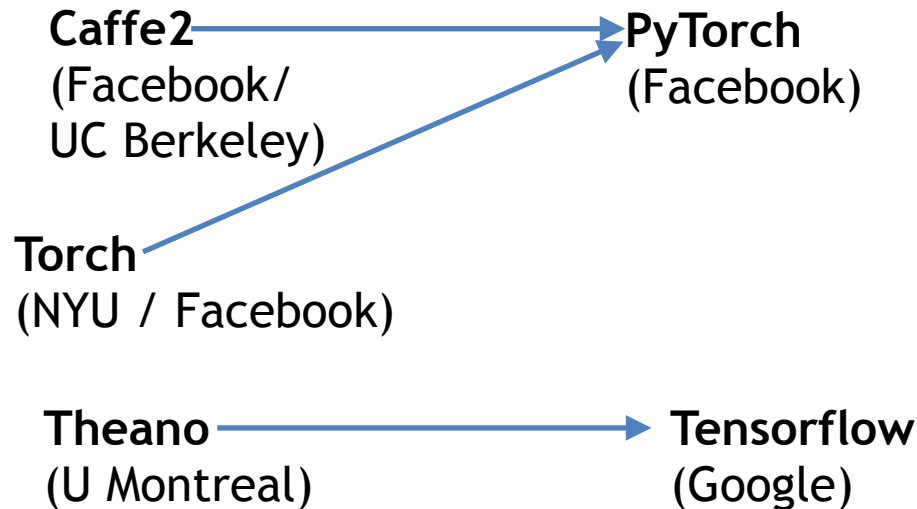
Torch

(NYU / Facebook)

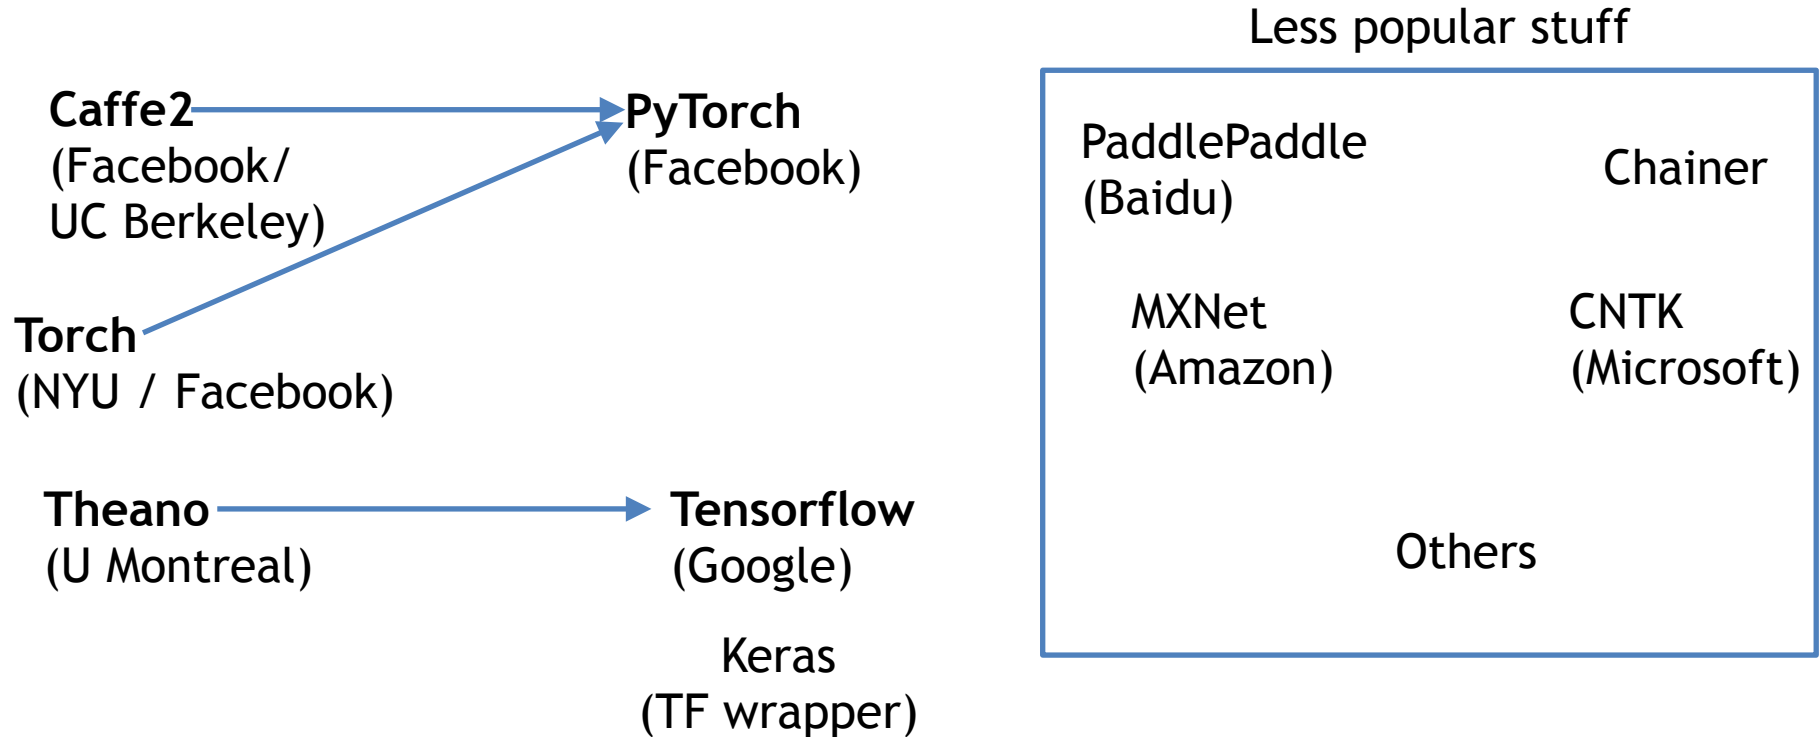
Theano

(U Montreal)

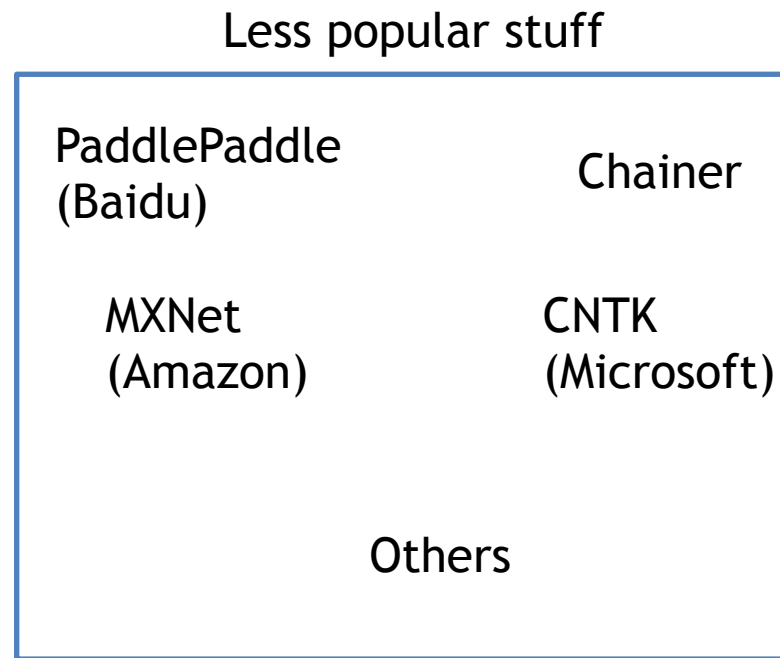
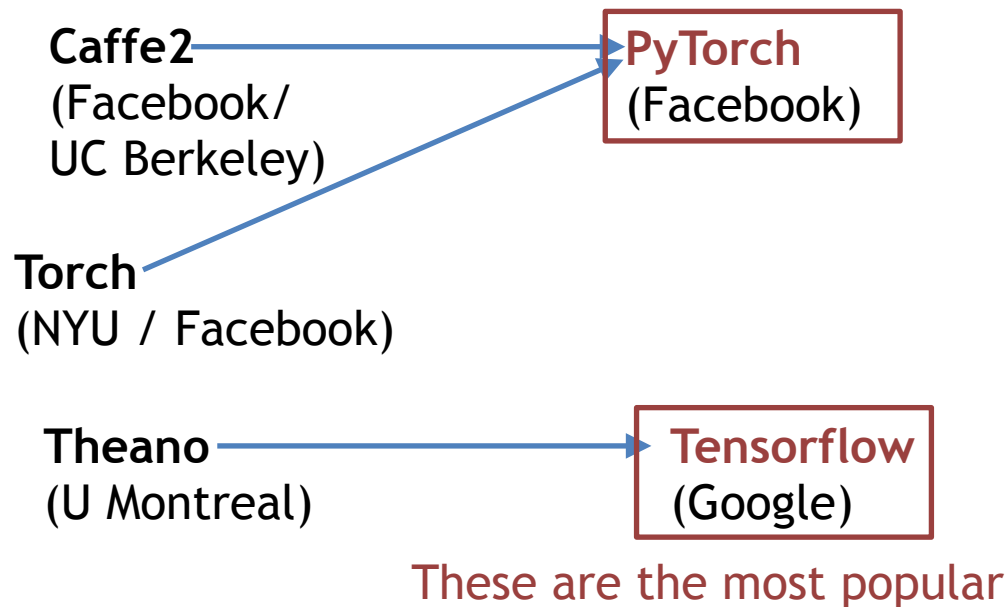
Frameworks



Frameworks




Frameworks



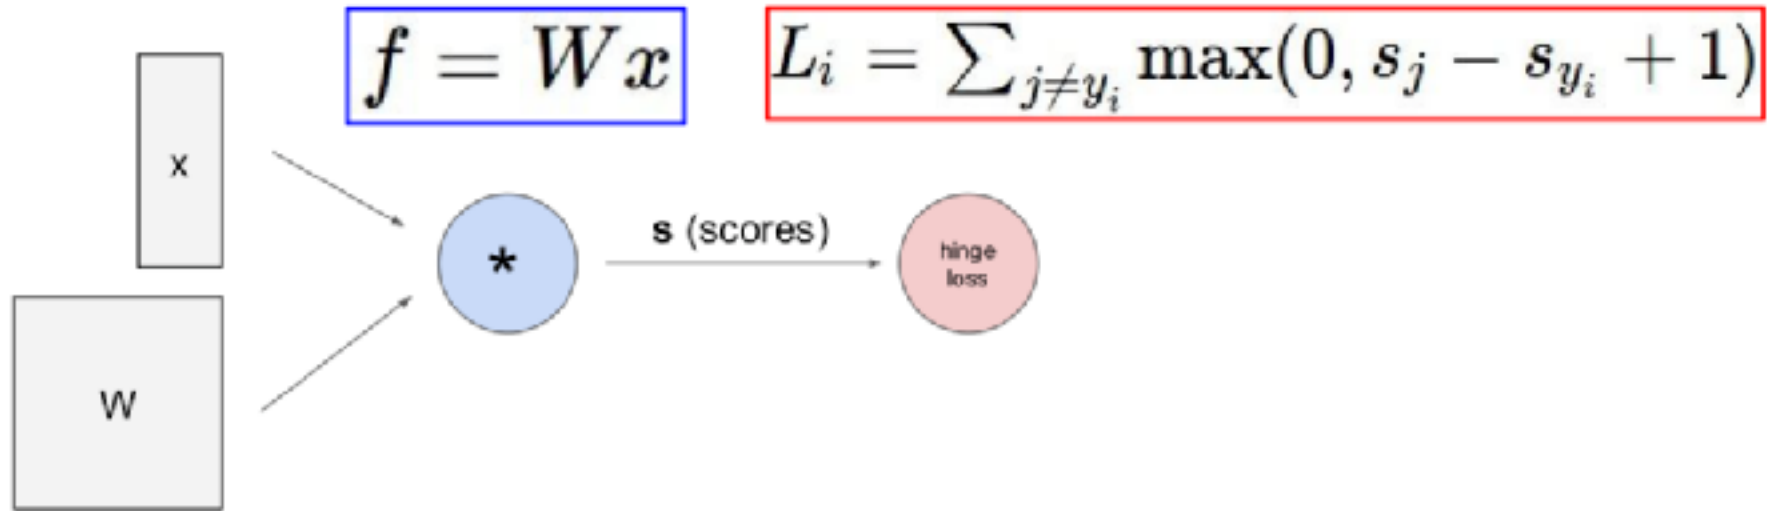
Neural Networks = Directed Acyclic Graphs

$$f = Wx \quad L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$

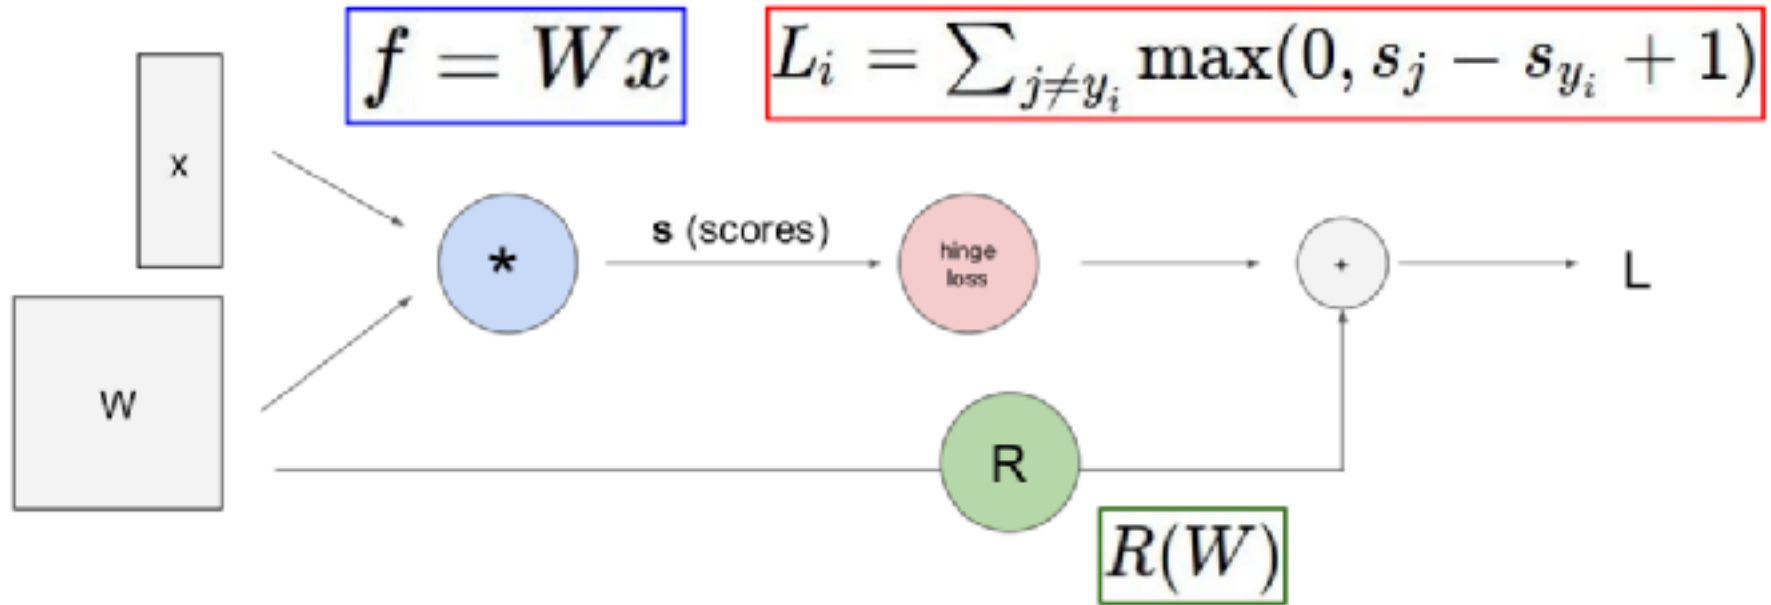
Neural Networks = Directed Acyclic Graphs




Neural Networks = Directed Acyclic Graphs



Neural Networks = Directed Acyclic Graphs



Computational Graph



Computational Graph


Numpy

```
import numpy as np
np.random.seed(0)

N, D = 3, 4

x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)

a = x * y
b = a + z
c = np.sum(b)
```



Computational Graph

Numpy


```
import numpy as np
np.random.seed(0)

N, D = 3, 4

x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)

a = x * y
b = a + z
c = np.sum(b)

grad_c = 1.0
grad_b = grad_c * np.ones((N, D))
grad_a = grad_b.copy()
grad_z = grad_b.copy()
grad_x = grad_a * y
grad_y = grad_a * x
```



Computational Graph

Numpy


```
import numpy as np
np.random.seed(0)

N, D = 3, 4

x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)

a = x * y
b = a + z
c = np.sum(b)

grad_c = 1.0
grad_b = grad_c * np.ones((N, D))
grad_a = grad_b.copy()
grad_z = grad_b.copy()
grad_x = grad_a * y
grad_y = grad_a * x
```



Good:

- Simple, clean API

Bad:

- Have to compute gradients ourself
- Can't run on GPU

Computational Graph

Numpy


```
import numpy as np
np.random.seed(0)

N, D = 3, 4

x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)

a = x * y
b = a + z
c = np.sum(b)

grad_c = 1.0
grad_b = grad_c * np.ones((N, D))
grad_a = grad_b.copy()
grad_z = grad_b.copy()
grad_x = grad_a * y
grad_y = grad_a * x
```



PyTorch

```
import torch
N, D = 3, 4
x = torch.randn(N, D)
y = torch.randn(N, D)
z = torch.randn(N, D)
a = x * y
b = a + z
c = torch.sum(b)
```

Forward Pass

Computational Graph

Numpy


```
import numpy as np
np.random.seed(0)

N, D = 3, 4

x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)

a = x * y
b = a + z
c = np.sum(b)

grad_c = 1.0
grad_b = grad_c * np.ones((N, D))
grad_a = grad_b.copy()
grad_z = grad_b.copy()
grad_x = grad_a * y
grad_y = grad_a * x
```



PyTorch

```
import torch
N, D = 3, 4
x = torch.randn(N, D, requires_grad=True)
y = torch.randn(N, D)
z = torch.randn(N, D)
a = x * y
b = a + z
c = torch.sum(b)
print(x.grad)
c.backward()
print(x.grad)
```


Computational Graph

Numpy


```
import numpy as np
np.random.seed(0)

N, D = 3, 4

x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)

a = x * y
b = a + z
c = np.sum(b)

grad_c = 1.0
grad_b = grad_c * np.ones((N, D))
grad_a = grad_b.copy()
grad_z = grad_b.copy()
grad_x = grad_a * y
grad_y = grad_a * x
```



PyTorch

```
import torch
N, D = 3, 4
x = torch.randn(N, D, requires_grad=True)
y = torch.randn(N, D)
z = torch.randn(N, D)
a = x * y
b = a + z
c = torch.sum(b)
print(x.grad)
c.backward()
print(x.grad)
```

None

tensor([[1.9707, 1.5089, 1.5041, -0.5507],
 [-0.1295, -0.6893, -0.2087, -0.0419],
 [0.0277, 1.4500, 1.3814, 0.3493]])

PyTorch (in-depth)

Pytorch: Fundamental Concepts

- Tensor: Like a numpy array, but can run on GPUs
- Module: A neural network layer; stores states and learnable weights

Images in Pytorch

In numpy, an image has shape:

- [Number of images, height, width, #colors]

In Pytorch, an image has shape:

- [Number of images, #colors, height, width]

Pytorch: Tensors

Example: A 2-layer neural network

```
1 I = 785
2 J = 64
3 C = 10
4 learning_rate = 1e-5
5 w1 = torch.randn(1, J, requires_grad=True)
6 w2 = torch.randn(J, C, requires_grad=True)
7
8 loss_function = torch.nn.NLLLoss()
9 losses = []
10 for i in range(10):
11     for (X,Y) in dataloader:
12         X = pre_process_image(X)
13         # forward pass
14         z_j = X.mm(w1)
15         a_j = torch.sigmoid(z_j)
16         z_k = a_j.mm(w2)
17         y_k = torch.softmax(z_k, dim=1)
18         # Compute loss
19         loss = loss_function(y_k, Y)
20         losses.append(loss)
21         # Backpropagation
22         loss.backward()
23         with torch.no_grad():
24             w1 -= learning_rate * w1.grad
25             w2 -= learning_rate * w2.grad
26             w1.grad.zero_()
27             w2.grad.zero_()
```

Pytorch: Tensors

Initialize the weights randomly

```
1 I = 785
2 J = 64
3 C = 10
4 learning_rate = 1e-5
5 w1 = torch.randn(I, J, requires_grad=True)
6 w2 = torch.randn(J, C, requires_grad=True)
7
8 loss_function = torch.nn.NLLLoss()
9 losses = []
10 for i in range(10):
11     for (X,Y) in dataloader:
12         X = pre_process_image(X)
13         # forward pass
14         z_j = X.mm(w1)
15         a_j = torch.sigmoid(z_j)
16         z_k = a_j.mm(w2)
17         y_k = torch.softmax(z_k, dim=1)
18         # Compute loss
19         loss = loss_function(y_k, Y)
20         losses.append(loss)
21         # Backpropagation
22         loss.backward()
23         with torch.no_grad():
24             w1 -= learning_rate * w1.grad
25             w2 -= learning_rate * w2.grad
26             w1.grad.zero_()
27             w2.grad.zero_()
```

Pytorch: Tensors

Do want to save gradients w.r.t weights

```
1 I = 785
2 J = 64
3 C = 10
4 learning_rate = 1e-5
5 w1 = torch.randn(1, J, requires_grad=True)
6 w2 = torch.randn(J, C, requires_grad=True)
7
8 loss_function = torch.nn.NLLLoss()
9 losses = []
10 for i in range(10):
11     for (X,Y) in dataloader:
12         X = pre_process_image(X)
13         # forward pass
14         z_j = X.mm(w1)
15         a_j = torch.sigmoid(z_j)
16         z_k = a_j.mm(w2)
17         y_k = torch.softmax(z_k, dim=1)
18         # Compute loss
19         loss = loss_function(y_k, Y)
20         losses.append(loss)
21         # Backpropagation
22         loss.backward()
23         with torch.no_grad():
24             w1 -= learning_rate * w1.grad
25             w2 -= learning_rate * w2.grad
26             w1.grad.zero_()
27             w2.grad.zero_()
```

Pytorch: Tensors

Define our loss function:
Cross Entropy Loss

```
1 T = 785
2 J = 64
3 C = 10
4 learning_rate = 1e-5
5 w1 = torch.randn(1, J, requires_grad=True)
6 w2 = torch.randn(J, C, requires_grad=True)
7
8 loss_function = torch.nn.NLLLoss()
9 losses = []
10 for i in range(10):
11     for (X,Y) in dataloader:
12         X = pre_process_image(X)
13         # forward pass
14         z_j = X.mm(w1)
15         a_j = torch.sigmoid(z_j)
16         z_k = a_j.mm(w2)
17         y_k = torch.softmax(z_k, dim=1)
18         # Compute loss
19         loss = loss_function(y_k, Y)
20         losses.append(loss)
21         # Backpropagation
22         loss.backward()
23     with torch.no_grad():
24         w1 -= learning_rate * w1.grad
25         w2 -= learning_rate * w2.grad
26         w1.grad.zero_()
27         w2.grad.zero_()
```


Pytorch: Tensors

Pytorch uses “dataloaders” to efficiently load datasets

```
1 I = 785
2 J = 64
3 C = 10
4 learning_rate = 1e-5
5 w1 = torch.randn(1, J, requires_grad=True)
6 w2 = torch.randn(J, C, requires_grad=True)
7
8 loss_function = torch.nn.NLLLoss()
9 losses = []
10 for i in range(10):
11     for (X,Y) in dataloader:
12         X = pre_process_image(X)
13         # forward pass
14         z_j = X.mm(w1)
15         a_j = torch.sigmoid(z_j)
16         z_k = a_j.mm(w2)
17         y_k = torch.softmax(z_k, dim=1)
18         # Compute loss
19         loss = loss_function(y_k, Y)
20         losses.append(loss)
21         # Backpropagation
22         loss.backward()
23         with torch.no_grad():
24             w1 -= learning_rate * w1.grad
25             w2 -= learning_rate * w2.grad
26             w1.grad.zero_()
27             w2.grad.zero_()
```

Pytorch: Tensors

Bias trick and normalization

```
1 I = 785
2 J = 64
3 C = 10
4 learning_rate = 1e-5
5 w1 = torch.randn(1, J, requires_grad=True)
6 w2 = torch.randn(J, C, requires_grad=True)
7
8 loss_function = torch.nn.NLLLoss()
9 losses = []
10 for i in range(10):
11     for (X,Y) in dataloader:
12         X = pre_process_image(X)
13         # forward pass
14         z_j = X.mm(w1)
15         a_j = torch.sigmoid(z_j)
16         z_k = a_j.mm(w2)
17         y_k = torch.softmax(z_k, dim=1)
18         # Compute loss
19         loss = loss_function(y_k, Y)
20         losses.append(loss)
21         # Backpropagation
22         loss.backward()
23         with torch.no_grad():
24             w1 -= learning_rate * w1.grad
25             w2 -= learning_rate * w2.grad
26             w1.grad.zero_()
27             w2.grad.zero_()
```

Pytorch: Tensors

Define the forward pass

```
1 I = 785
2 J = 64
3 C = 10
4 learning_rate = 1e-5
5 w1 = torch.randn(1, J, requires_grad=True)
6 w2 = torch.randn(J, C, requires_grad=True)
7
8 loss_function = torch.nn.NLLLoss()
9 losses = []
10 for i in range(10):
11     for (X,Y) in dataloader:
12         X = pre_process_image(X)
13         # forward pass
14         z_j = X.mm(w1)
15         a_j = torch.sigmoid(z_j)
16         z_k = a_j.mm(w2)
17         y_k = torch.softmax(z_k, dim=1)
18         # Compute loss
19         loss = loss_function(y_k, Y)
20         losses.append(loss)
21         # Backpropagation
22         loss.backward()
23         with torch.no_grad():
24             w1 -= learning_rate * w1.grad
25             w2 -= learning_rate * w2.grad
26             w1.grad.zero_()
27             w2.grad.zero_()
```

Pytorch: Tensors

Compute the loss



```
1 I = 785
2 J = 64
3 C = 10
4 learning_rate = 1e-5
5 w1 = torch.randn(1, J, requires_grad=True)
6 w2 = torch.randn(J, C, requires_grad=True)
7
8 loss_function = torch.nn.NLLLoss()
9 losses = []
10 for i in range(10):
11     for (X,Y) in dataloader:
12         X = pre_process_image(X)
13         # forward pass
14         z_j = X.mm(w1)
15         a_j = torch.sigmoid(z_j)
16         z_k = a_j.mm(w2)
17         y_k = torch.softmax(z_k, dim=1)
18         # Compute loss
19         loss = loss_function(y_k, Y)
20         losses.append(loss)
21         # Backpropagation
22         loss.backward()
23         with torch.no_grad():
24             w1 -= learning_rate * w1.grad
25             w2 -= learning_rate * w2.grad
26             w1.grad.zero_()
27             w2.grad.zero_()
```

Pytorch: Tensors

Backpropagate the loss



```
1 I = 785
2 J = 64
3 C = 10
4 learning_rate = 1e-5
5 w1 = torch.randn(1, J, requires_grad=True)
6 w2 = torch.randn(J, C, requires_grad=True)
7
8 loss_function = torch.nn.NLLLoss()
9 losses = []
10 for i in range(10):
11     for (X,Y) in dataloader:
12         X = pre_process_image(X)
13         # forward pass
14         z_j = X.mm(w1)
15         a_j = torch.sigmoid(z_j)
16         z_k = a_j.mm(w2)
17         y_k = torch.softmax(z_k, dim=1)
18         # Compute loss
19         loss = loss_function(y_k, Y)
20         losses.append(loss)
21         # Backpropagation
22         loss.backward()
23     with torch.no_grad():
24         w1 -= learning_rate * w1.grad
25         w2 -= learning_rate * w2.grad
26         w1.grad.zero_()
27         w2.grad.zero_()
```

Pytorch: Tensors

Make gradient step on weights

torch.no_grad() means “don’t build a computational graph here”

```
1 I = 785
2 J = 64
3 C = 10
4 learning_rate = 1e-5
5 w1 = torch.randn(1, J, requires_grad=True)
6 w2 = torch.randn(J, C, requires_grad=True)
7
8 loss_function = torch.nn.NLLLoss()
9 losses = []
10 for i in range(10):
11     for (X,Y) in dataloader:
12         X = pre_process_image(X)
13         # forward pass
14         z_j = X.mm(w1)
15         a_j = torch.sigmoid(z_j)
16         z_k = a_j.mm(w2)
17         y_k = torch.softmax(z_k, dim=1)
18         # Compute loss
19         loss = loss_function(y_k, Y)
20         losses.append(loss)
21         # Backpropagation
22         loss.backward()
23         with torch.no_grad():
24             w1 -= learning_rate * w1.grad
25             w2 -= learning_rate * w2.grad
26             w1.grad.zero_()
27             w2.grad.zero_()
```

Pytorch: torch.nn

Higher lever wrapper for defining neural networks

```
1 I = 785
2 J = 64
3 C = 10
4 learning_rate = 1e-5
5 model = nn.Sequential(
6     nn.Linear(I, J),
7     nn.Sigmoid(),
8     nn.Linear(J, C)
9     # No need for softmax, since its included in
10    # nn.CrossEntropyLoss()
11 )
12
13 loss_function = torch.nn.CrossEntropyLoss()
14 losses = []
15 for epoch in range(10):
16     for [X,Y] in dataloader:
17         X = pre_process_image(X)
18         # forward pass
19         y_k = model(X)
20         # Compute loss
21         loss = loss_function(y_k, Y)
22         losses.append(loss)
23         # Backpropagation
24         loss.backward()
25         with torch.no_grad():
26             for param in model.parameters():
27                 param -= learning_rate * param.grad
```

Pytorch: torch.nn

Higher lever wrapper for defining neural networks

Define each layer in model.

Each layer is a nn.Module() object, containing learnable weights.

```
1 I = 785
2 J = 64
3 C = 10
4 learning_rate = 1e-5
5 model = nn.Sequential(
6     nn.Linear(I, J),
7     nn.Sigmoid(),
8     nn.Linear(J, C)
9     # No need for softmax, since its included in
10    # nn.CrossEntropyLoss()
11 )
12
13 loss_function = torch.nn.CrossEntropyLoss()
14 losses = []
15 for epoch in range(10):
16     for [X,Y] in dataloader:
17         X = pre_process_image(X)
18         # forward pass
19         y_k = model(X)
20         # Compute loss
21         loss = loss_function(y_k, Y)
22         losses.append(loss)
23         # Backpropagation
24         loss.backward()
25         with torch.no_grad():
26             for param in model.parameters():
27                 param -= learning_rate * param.grad
```


Pytorch: torch.nn

Changed loss function to
nn.CrossEntropyLoss.
This includes the softmax!

```
1 I = 785
2 J = 64
3 C = 10
4 learning_rate = 1e-5
5 model = nn.Sequential(
6     nn.Linear(I, J),
7     nn.Sigmoid(),
8     nn.Linear(J, C)
9     # No need for softmax, since its included in
10    # nn.CrossEntropyLoss()
11 )
12
13 loss_function = torch.nn.CrossEntropyLoss()
14 losses = []
15 for epoch in range(10):
16     for [X,Y] in dataloader:
17         X = pre_process_image(X)
18         # forward pass
19         y_k = model(X)
20         # Compute loss
21         loss = loss_function(y_k, Y)
22         losses.append(loss)
23         # Backpropagation
24         loss.backward()
25         with torch.no_grad():
26             for param in model.parameters():
27                 param -= learning_rate * param.grad
```

Pytorch: torch.nn

Simplifies our forward pass!

```
1 I = 785
2 J = 64
3 C = 10
4 learning_rate = 1e-5
5 model = nn.Sequential(
6     nn.Linear(I, J),
7     nn.Sigmoid(),
8     nn.Linear(J, C)
9     # No need for softmax, since its included in
10    # nn.CrossEntropyLoss()
11 )
12
13 loss_function = torch.nn.CrossEntropyLoss()
14 losses = []
15 for epoch in range(10):
16     for [X,Y] in dataloader:
17         X = pre_process_image(X)
18         # forward pass
19         y_k = model(X)
20         # Compute loss
21         loss = loss_function(y_k, Y)
22         losses.append(loss)
23         # Backpropagation
24         loss.backward()
25         with torch.no_grad():
26             for param in model.parameters():
27                 param -= learning_rate * param.grad
```

Pytorch: torch.nn

Compute loss and perform backward pass

Each weight in model has `requires_grad=True` by default

```
1 I = 785
2 J = 64
3 C = 10
4 learning_rate = 1e-5
5 model = nn.Sequential(
6     nn.Linear(I, J),
7     nn.Sigmoid(),
8     nn.Linear(J, C)
9     # No need for softmax, since its included in
10    # nn.CrossEntropyLoss()
11 )
12
13 loss_function = torch.nn.CrossEntropyLoss()
14 losses = []
15 for epoch in range(10):
16     for [X,Y] in dataloader:
17         X = pre_process_image(X)
18         # forward pass
19         y_k = model(X)
20         # Compute loss
21         loss = loss_function(y_k, Y)
22         losses.append(loss)
23         # Backpropagation
24         loss.backward()
25     with torch.no_grad():
26         for param in model.parameters():
27             param -= learning_rate * param.grad
```

Pytorch: torch.nn

Perform our gradient step
(and disable gradients)

```
1 I = 785
2 J = 64
3 C = 10
4 learning_rate = 1e-5
5 model = nn.Sequential(
6     nn.Linear(I, J),
7     nn.Sigmoid(),
8     nn.Linear(J, C)
9     # No need for softmax, since its included in
10    # nn.CrossEntropyLoss()
11 )
12
13 loss_function = torch.nn.CrossEntropyLoss()
14 losses = []
15 for epoch in range(10):
16     for [X,Y] in dataloader:
17         X = pre_process_image(X)
18         # forward pass
19         y_k = model(X)
20         # Compute loss
21         loss = loss_function(y_k, Y)
22         losses.append(loss)
23         # Backpropagation
24         loss.backward()
25         with torch.no_grad():
26             for param in model.parameters():
27                 param -= learning_rate * param.grad
```

Pytorch: torch.optim

Final piece you need to know

Implements **Stochastic Gradient Descent**

Input: our learnable parameters (weights + biases)
+ learning rate

```
1 I = 785
2 J = 64
3 C = 10
4 learning_rate = 1e-5
5 model = nn.Sequential(
6     nn.Linear(I, J),
7     nn.Sigmoid(),
8     nn.Linear(J, C)
9     # No need for softmax, since its included in
10    # nn.CrossEntropyLoss()
11 )
12
13 loss_function = torch.nn.CrossEntropyLoss()
14 optimizer = torch.optim.SGD(model.parameters(),
15                               lr=learning_rate)
16 losses = []
17 for epoch in range(10):
18     for (X,Y) in dataloader:
19         X = pre_process_image(X)
20         # forward pass
21         y_k = model(X)
22         # Compute loss
23         loss = loss_function(y_k, Y)
24         losses.append(loss)
25         # Backpropagation
26         loss.backward()
27         optimizer.step()
28         optimizer.zero_grad()
```

Pytorch: torch.optim

Perform gradient step and reset the gradients

```
1 I = 785
2 J = 64
3 C = 10
4 learning_rate = 1e-5
5 model = nn.Sequential(
6     nn.Linear(I, J),
7     nn.Sigmoid(),
8     nn.Linear(J, C)
9     # No need for softmax, since its included in
10    # nn.CrossEntropyLoss()
11 )
12
13 loss_function = torch.nn.CrossEntropyLoss()
14 optimizer = torch.optim.SGD(model.parameters(),
15                               lr=learning_rate)
16 losses = []
17 for epoch in range(10):
18     for (X,Y) in dataloader:
19         X = pre_process_image(X)
20         # forward pass
21         y_k = model(X)
22         # Compute loss
23         loss = loss_function(y_k, Y)
24         losses.append(loss)
25         # Backpropagation
26         loss.backward()
27         optimizer.step()
28         optimizer.zero_grad()
```

Pytorch: nn.Module

A PyTorch **Module** is a neural network layer; it inputs and outputs tensors

Can contain weights or other modules

Required for more complex layers

Easily customizable layers

```
1 class TwoLayerNet(nn.Module):
2     def __init__(self):
3         super().__init__()
4         I, J, C = 785, 64, 10
5         self.layer1 = nn.Sequential(
6             nn.Linear(I, J),
7             nn.Sigmoid()
8         )
9         self.layer2 = nn.Linear(J, C)
10    def forward(self, x):
11        z = self.layer1(x)
12        x = self.layer2(z)
13        return x
14
15    learning_rate = 1e-3
16    model = TwoLayerNet()
17
18    loss_function = torch.nn.CrossEntropyLoss()
19    optimizer = torch.optim.SGD(model.parameters(),
20                                lr=learning_rate)
21    losses = []
22    for epoch in range(2):
23        for (X, Y) in dataloader:
24            X = pre_process_image(X)
25            # forward pass
26            y_k = model(X)
27            # Compute loss
28            loss = loss_function(y_k, Y)
29            losses.append(loss)
30            # Backpropagation
31            loss.backward()
32            optimizer.step()
33            optimizer.zero_grad()
34
```

Pytorch: nn.Module

Start with defining the model

```
1 class TwoLayerNet(nn.Module):
2     def __init__(self):
3         super().__init__()
4         I, J, C = 785, 64, 10
5         self.layer1 = nn.Sequential(
6             nn.Linear(I, J),
7             nn.Sigmoid()
8         )
9         self.layer2 = nn.Linear(J, C)
10    def forward(self, x):
11        z = self.layer1(x)
12        x = self.layer2(x)
13        return x
14
15    learning_rate = 1e-3
16    model = TwoLayerNet()
17
18    loss_function = torch.nn.CrossEntropyLoss()
19    optimizer = torch.optim.SGD(model.parameters(),
20                                lr=learning_rate)
21    losses = []
22    for epoch in range(2):
23        for (X,Y) in dataloader:
24            X = pre_process_image(X)
25            # forward pass
26            y_k = model(X)
27            # Compute loss
28            loss = loss_function(y_k, Y)
29            losses.append(loss)
30            # Backpropagation
31            loss.backward()
32            optimizer.step()
33            optimizer.zero_grad()
34
```


Pytorch: nn.Module

Called when we initialize our model

```
1 class TwoLayerNet(nn.Module):
2     def __init__(self):
3         super().__init__()
4         I, J, C = 785, 64, 10
5         self.layer1 = nn.Sequential(
6             nn.Linear(I, J),
7             nn.Sigmoid()
8         )
9         self.layer2 = nn.Linear(J, C)
10    def forward(self, x):
11        z = self.layer1(x)
12        x = self.layer2(z)
13        return x
14
15    learning_rate = 1e-3
16    model = TwoLayerNet()
17
18    loss_function = torch.nn.CrossEntropyLoss()
19    optimizer = torch.optim.SGD(model.parameters(),
20                                lr=learning_rate)
21    losses = []
22    for epoch in range(2):
23        for (X, Y) in dataloader:
24            X = pre_process_image(X)
25            # forward pass
26            y_k = model(X)
27            # Compute loss
28            loss = loss_function(y_k, Y)
29            losses.append(loss)
30            # Backpropagation
31            loss.backward()
32            optimizer.step()
33            optimizer.zero_grad()
34
```

Pytorch: nn.Module

Called when we perform forward pass

```
1 class TwoLayerNet(nn.Module):
2     def __init__(self):
3         super().__init__()
4         I, J, C = 785, 64, 10
5         self.layer1 = nn.Sequential(
6             nn.Linear(I, J),
7             nn.Sigmoid()
8         )
9         self.layer2 = nn.Linear(J, C)
10    def forward(self, x):
11        x = self.layer1(x)
12        x = self.layer2(x)
13        return x
14
15    learning_rate = 1e-3
16    model = TwoLayerNet()
17
18    loss_function = torch.nn.CrossEntropyLoss()
19    optimizer = torch.optim.SGD(model.parameters(),
20                                lr=learning_rate)
21    losses = []
22    for epoch in range(2):
23        for (X, Y) in dataloader:
24            X = pre_process_image(X)
25            # forward pass
26            y_k = model(X)
27            # Compute loss
28            loss = loss_function(y_k, Y)
29            losses.append(loss)
30            # Backpropagation
31            loss.backward()
32            optimizer.step()
33            optimizer.zero_grad()
34
```

Pytorch: DataLoaders

A **DataLoader** wraps a dataset and provides features such as:

- Data augmentation
- Data pre-processing
- mini-batch shuffling and splitting

```
1 class TwoLayerNet(nn.Module):
2     def __init__(self):
3         super().__init__()
4         I, J, C = 784, 64, 10
5         self.layer1 = nn.Sequential(
6             nn.Linear(I, J),
7             nn.Sigmoid()
8         )
9         self.layer2 = nn.Linear(J, C)
10    def forward(self, x):
11        x = self.layer1(x)
12        x = self.layer2(x)
13        return x
14
15    learning_rate = 1e-3
16    batch_size=32
17    model = TwoLayerNet()
18
19    dataloader_train, dataloader_test = load_mnist(batch_size
20
21    loss_function = torch.nn.CrossEntropyLoss()
22    optimizer = torch.optim.SGD(model.parameters(),
23                                  lr=learning_rate)
24    losses = []
25    for epoch in range(2):
26        for (X_batch, Y_batch) in dataloader_train:
27            X_batch = pre_process_image(X_batch)
28            # forward pass
29            y_k = model(X_batch)
30            # Compute loss
31            loss = loss_function(y_k, Y_batch)
32            losses.append(loss)
33            # Backpropagation
34            loss.backward()
35            optimizer.step()
36            optimizer.zero_grad()
```

Pytorch: DataLoaders

Iterates over each batch in a epoch



```
1 class TwoLayerNet(nn.Module):
2     def __init__(self):
3         super().__init__()
4         I, J, C = 784, 64, 10
5         self.layer1 = nn.Sequential(
6             nn.Linear(I, J),
7             nn.Sigmoid()
8         )
9         self.layer2 = nn.Linear(J, C)
10    def forward(self, x):
11        x = self.layer1(x)
12        x = self.layer2(x)
13        return x
14
15    learning_rate = 1e-3
16    batch_size=32
17    model = TwoLayerNet()
18
19    dataloader_train, dataloader_test = load_mnist(batch_size
20
21    loss_function = torch.nn.CrossEntropyLoss()
22    optimizer = torch.optim.SGD(model.parameters(),
23                                lr=learning_rate)
24
25    losses = []
26    for epoch in range(2):
27        for (X_batch, Y_batch) in dataloader_train:
28            X_batch = pre_process_image(X_batch)
29            # forward pass
30            y_k = model(X_batch)
31            # Compute loss
32            loss = loss_function(y_k, Y_batch)
33            losses.append(loss)
34            # Backpropagation
35            loss.backward()
36            optimizer.step()
37            optimizer.zero_grad()
```

Deep Learning Hardware

CPU, GPU, TPU

We have two hardware choices:

- NVIDIA GPU
- Google Tensor Processing Unit (TPU)
- AMD? Really not used.


GPU vs CPU

	Cores	Clock Speed	Memory	Price	Speed
CPU (Intel Core i7-7700k)	4 (8 threads with hyperthreading)	4.2 GHz	System RAM	\$339	~540 GFLOPs FP32
GPU (NVIDIA GTX 1080 Ti)	3584	1.6 GHz	11 GB GDDR5 X	\$699	~11.4 TFLOPs FP32

GPU: thousands of “dumb” cores: Great for parallel tasks

Neural Networks: “Only” matrix multiplication, easy in parallel

GPU vs CPU for CNNs



GPU vs CPU: ResNet-200

- Forward pass:
 - Pascal Titan X: 104ms
 - CPU: Dual Xeon E5-2630 v3: 8,666ms (83x slower)
- Backward pass:
 - Pascal Titan X: 191 ms
 - CPU: Dual Xeon E5-2630 v3: 13,758 ms (72x slower)

Pytorch: On GPU

Utilizing GPU resources is simple!

Note, it is not required for assignment 1.
Recommended for assignment 2!

```
1 class TwoLayerNet(nn.Module):
2     def __init__(self):
3         super().__init__()
4         I, J, C = 784, 64, 10
5         self.layer1 = nn.Sequential(
6             nn.Linear(I, J),
7             nn.Sigmoid()
8         )
9         self.layer2 = nn.Linear(J, C)
10    def forward(self, x):
11        x = self.layer1(x)
12        x = self.layer2(x)
13        return x
14
15    learning_rate = 1e-3
16    batch_size=32
17    model = TwoLayerNet().cuda()
18
19    dataloader_train, dataloader_test = load_mnist(batch_size)
20
21    loss_function = torch.nn.CrossEntropyLoss()
22    optimizer = torch.optim.SGD(model.parameters(),
23                                  lr=learning_rate)
24    losses = []
25    for epoch in range(2):
26        for (X_batch, Y_batch) in dataloader_train:
27            X_batch = pre_process_image(X_batch)
28            X_batch, Y_batch = X_batch.cuda(), Y_batch.cuda()
29            # forward pass
30            y_k = model(X_batch)
31            # Compute loss
32            loss = loss_function(y_k, Y_batch)
33            losses.append(loss)
34            # Backpropagation
35            loss.backward()
36            optimizer.step()
37            optimizer.zero_grad()
```

Pytorch: On GPU

Utilizing GPU resources is simple!

`.cuda()` transfers weights/tensors to GPU VRAM

```
1 class TwoLayerNet(nn.Module):
2     def __init__(self):
3         super().__init__()
4         I, J, C = 784, 64, 10
5         self.layer1 = nn.Sequential(
6             nn.Linear(I, J),
7             nn.Sigmoid()
8         )
9         self.layer2 = nn.Linear(J, C)
10    def forward(self, x):
11        x = self.layer1(x)
12        x = self.layer2(x)
13        return x
14
15    learning_rate = 1e-3
16    batch_size=32
17    model = TwoLayerNet().cuda()
18
19    dataloader_train, dataloader_test = load_mnist(batch_size=
20
21    loss_function = torch.nn.CrossEntropyLoss()
22    optimizer = torch.optim.SGD(model.parameters(),
23                                  lr=learning_rate)
24    losses = []
25    for epoch in range(2):
26        for (X_batch, Y_batch) in dataloader_train:
27            X_batch = pre_process_image(X_batch)
28            X_batch, Y_batch = X_batch.cuda(), Y_batch.cuda()
29            # forward pass
30            y_k = model(X_batch)
31            # Compute loss
32            loss = loss_function(y_k, Y_batch)
33            losses.append(loss)
34            # Backpropagation
35            loss.backward()
36            optimizer.step()
37            optimizer.zero_grad()
```

Pytorch: On GPU

Utilizing GPU resources is simple!

`.cuda()` transfers weights/tensors to GPU VRAM

```
1 class TwoLayerNet(nn.Module):
2     def __init__(self):
3         super().__init__()
4         I, J, C = 784, 64, 10
5         self.layer1 = nn.Sequential(
6             nn.Linear(I, J),
7             nn.Sigmoid()
8         )
9         self.layer2 = nn.Linear(J, C)
10    def forward(self, x):
11        x = self.layer1(x)
12        x = self.layer2(x)
13        return x
14
15    learning_rate = 1e-3
16    batch_size=32
17    model = TwoLayerNet().cuda()
18
19    dataloader_train, dataloader_test = load_mnist(batch_size)
20
21    loss_function = torch.nn.CrossEntropyLoss()
22    optimizer = torch.optim.SGD(model.parameters(),
23                                  lr=learning_rate)
24    losses = []
25    for epoch in range(2):
26        for (X_batch, Y_batch) in dataloader_train:
27            X_batch = pre_process_image(X_batch)
28            X_batch, Y_batch = X_batch.cuda(), Y_batch.cuda()
29            # forward pass
30            y_k = model(X_batch)
31            # Compute loss
32            loss = loss_function(y_k, Y_batch)
33            losses.append(loss)
34            # Backpropagation
35            loss.backward()
36            optimizer.step()
37            optimizer.zero_grad()
```

Pytorch: On GPU

Utilizing GPU resources is simple!

`.cuda()` transfers weights/tensors to GPU VRAM

CAREFUL: Calling `.cuda()` without a NVIDIA GPU available will cause error!

```
1 class TwoLayerNet(nn.Module):
2     def __init__(self):
3         super().__init__()
4         I, J, C = 784, 64, 10
5         self.layer1 = nn.Sequential(
6             nn.Linear(I, J),
7             nn.Sigmoid()
8         )
9         self.layer2 = nn.Linear(J, C)
10    def forward(self, x):
11        x = self.layer1(x)
12        x = self.layer2(x)
13        return x
14
15    learning_rate = 1e-3
16    batch_size=32
17    model = TwoLayerNet().cuda()
18
19    dataloader_train, dataloader_test = load_mnist(batch_size=
20
21    loss_function = torch.nn.CrossEntropyLoss()
22    optimizer = torch.optim.SGD(model.parameters(),
23                                  lr=learning_rate)
24    losses = []
25    for epoch in range(2):
26        for (X_batch, Y_batch) in dataloader_train:
27            X_batch = pre_process_image(X_batch)
28            X_batch, Y_batch = X_batch.cuda(), Y_batch.cuda()
29            # forward pass
30            y_k = model(X_batch)
31            # Compute loss
32            loss = loss_function(y_k, Y_batch)
33            losses.append(loss)
34            # Backpropagation
35            loss.backward()
36            optimizer.step()
37            optimizer.zero_grad()
```

Pytorch: On GPU

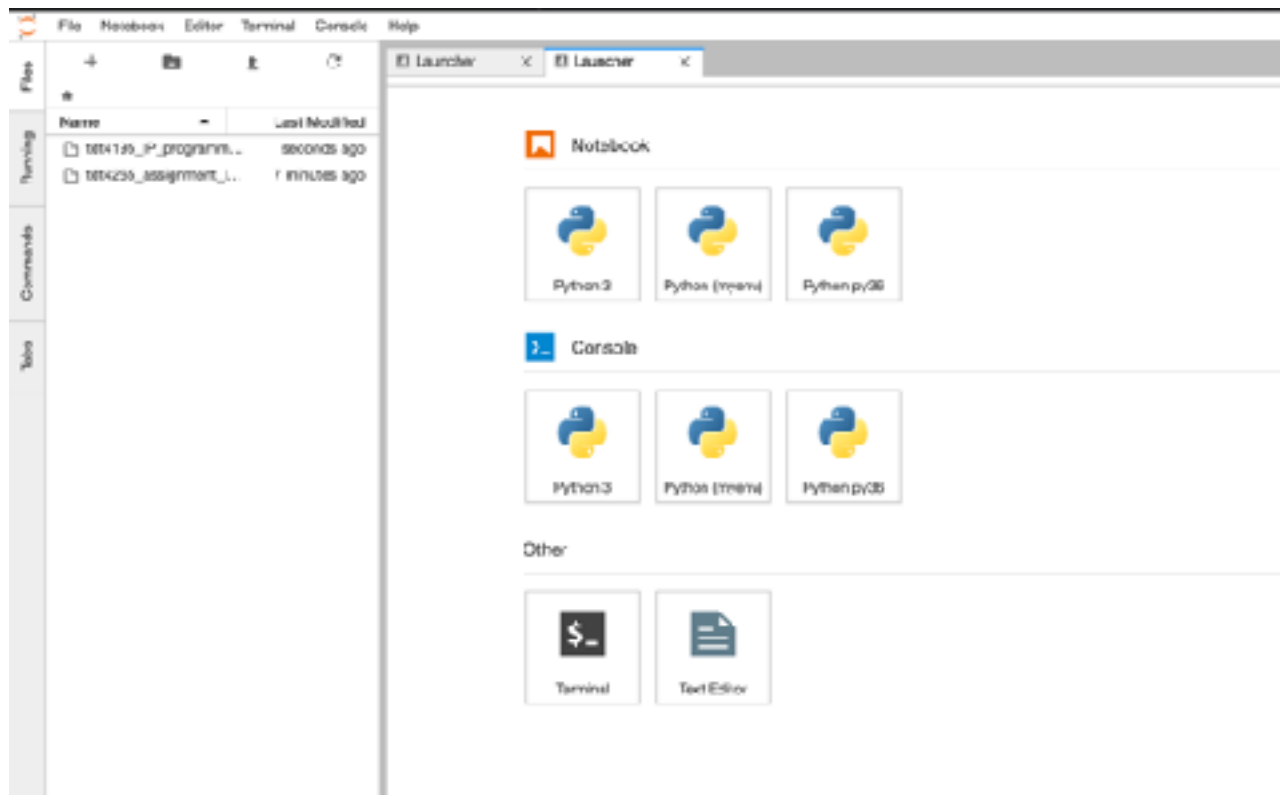
Instead: Implement a to_cuda() function

```
1 def to_cuda(elements):
2     if torch.cuda.is_available():
3         if type(elements) == tuple or type(elements) == list:
4             return [x.cuda() for x in elements]
5         return elements.cuda()
6     return elements
```

```
1 class TwoLayerNet(nn.Module):
2     def __init__(self):
3         super().__init__()
4         I, J, C = 784, 64, 10
5         self.layer1 = nn.Sequential(
6             nn.Linear(I, J),
7             nn.Sigmoid()
8         )
9         self.layer2 = nn.Linear(J, C)
10    def forward(self, x):
11        x = self.layer1(x)
12        x = self.layer2(x)
13        return x
14
15    learning_rate = 1e-3
16    batch_size=32
17    model = to_cuda(TwoLayerNet())
18
19    dataloader_train, dataloader_test = load_mnist(batch_size)
20
21    loss_function = torch.nn.CrossEntropyLoss()
22    optimizer = torch.optim.SGD(model.parameters(),
23                                  lr=learning_rate)
24    losses = []
25    for epoch in range(2):
26        for (X_batch, Y_batch) in dataloader_train:
27            X_batch = pre_process_image(X_batch)
28            X_batch, Y_batch = to_cuda([X_batch, Y_batch])
29            # forward pass
30            y_k = model(X_batch)
31            # Compute loss
32            loss = loss_function(y_k, Y_batch)
33            losses.append(loss)
34            # Backpropagation
35            loss.backward()
36            optimizer.step()
37            optimizer.zero_grad()
```

Cool features & resources

Jupyter Notebook/Lab



Multi-layer neural networks

- What can a 1-layer network predict?
<https://bit.ly/2HO6iln>

Multi-layer neural networks

- What can a 1-layer network predict?
<https://bit.ly/2HO6iln>
- Answer: Only linearly separable functions

Multi-layer neural networks

- What can a 2-layer network predict? (1 hidden layer)
<https://bit.ly/2SndBEQ>

Multi-layer neural networks

- What can a 2-layer network predict? (1 hidden layer)
<https://bit.ly/2SndBEQ>
- Answer: Theoretically, everything possible.

Pytorch model zoo

Pytorch comes with a large amount of state-of-the-art models

Classification:

- <https://pytorch.org/docs/stable/torchvision/models.html>

Detection:

- <https://pytorch.org/blog/torchvision03/>

Google Colab

Google Colab is a jupyter notebook like system
With **FREE** GPU resources!

Some cool colabs:

- [Detectron2](#)
- [DeepPrivay \(Shameless plug\)](#)
- [BigGAN](#)

Pytorch model zoo

```
1 # Define model
2 model = torchvision.models.resnet152(pretrained=True)
3 model = model.eval()
4
5 im = skimage.data.chelsea()
6 im = Image.fromarray(im)
7 plt.imshow(im)
8
9 im = transform(im)
10 preds = model(im[None])
11 print(get_imageNet_class(preds))
```

Pytorch model zoo

tiger_cat

